


CS330: Operating Systems

Virtual memory: Address spaces


Recap: The process abstraction

- The OS creates a *process* when we run an *executable*



- Executable is a file, stored in a persistent storage (e.g., disk)
- To run, the process code and data should reside in memory
- Run-time memory allocation and deallocation should be supported

Executable file to process memory view



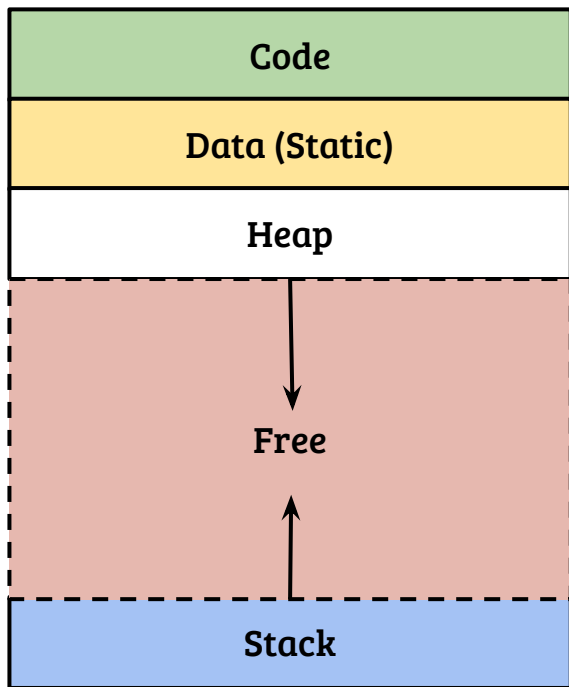
- A typical executable file contains code and statically allocated data
- Statically allocated: global and static variables
- Is loading the program (code and data) sufficient for program execution?

Executable file to process memory view




- A typical executable file contains code and statically allocated data
- Statically allocated: global and static variables
- Is loading the program (code and data) sufficient for program execution?
- No, we need memory for stack and dynamic allocation
- Stack: function call and return, store local (stack) variables
- Heap: dynamic memory allocation through APIs like *malloc()*

The address space abstraction




- Address space represents memory state of a process
- Address space layout is same for all the processes (convenience)
- Exact layout can be decided by the OS, conventional layout is shown

The address space abstraction



- If all processes have same address space, how they map to actual memory?
- What are the responsibilities of the OS during program load?
 - How CPU register state is changed?
- What is the OS role in dynamic memory allocation?




Exact layout can be decided by the OS,
conventional layout is shown

The address space abstraction

Code

- If all processes have same address space, how they map to actual memory?
- Architecture support used by OS techniques to perform memory virtualization i.e., translate virtual address to physical address (will revisit)
- What are the responsibilities of the OS during program load?
 - How CPU register state is changed?
- What is the OS role in dynamic memory allocation?

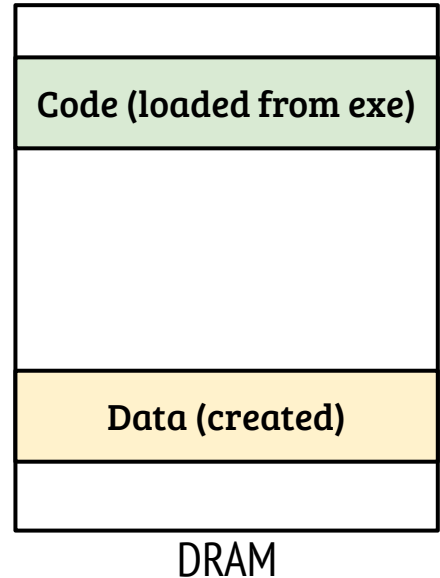

OS during program load (exec)



- A fresh address space is initialized
- In reality, parent address space copied at the time of `fork()` is reset
- PC and SP are set with addresses of code and stack, respectively

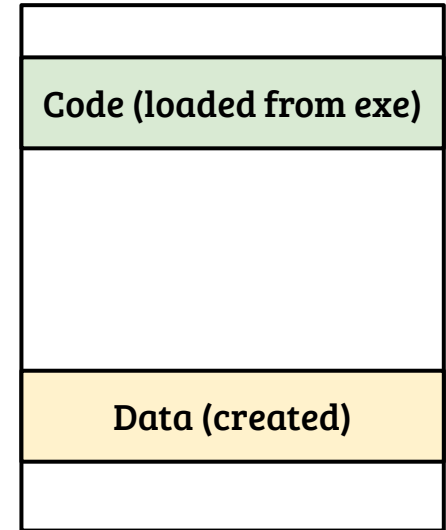

OS during program load (exec)

- Physical memory for code and data allocated, executable code (text section) is loaded



OS during program load (exec)

- Translation information updated
- Process is ready to execute
- Executes when register state in PCB is loaded onto the CPU

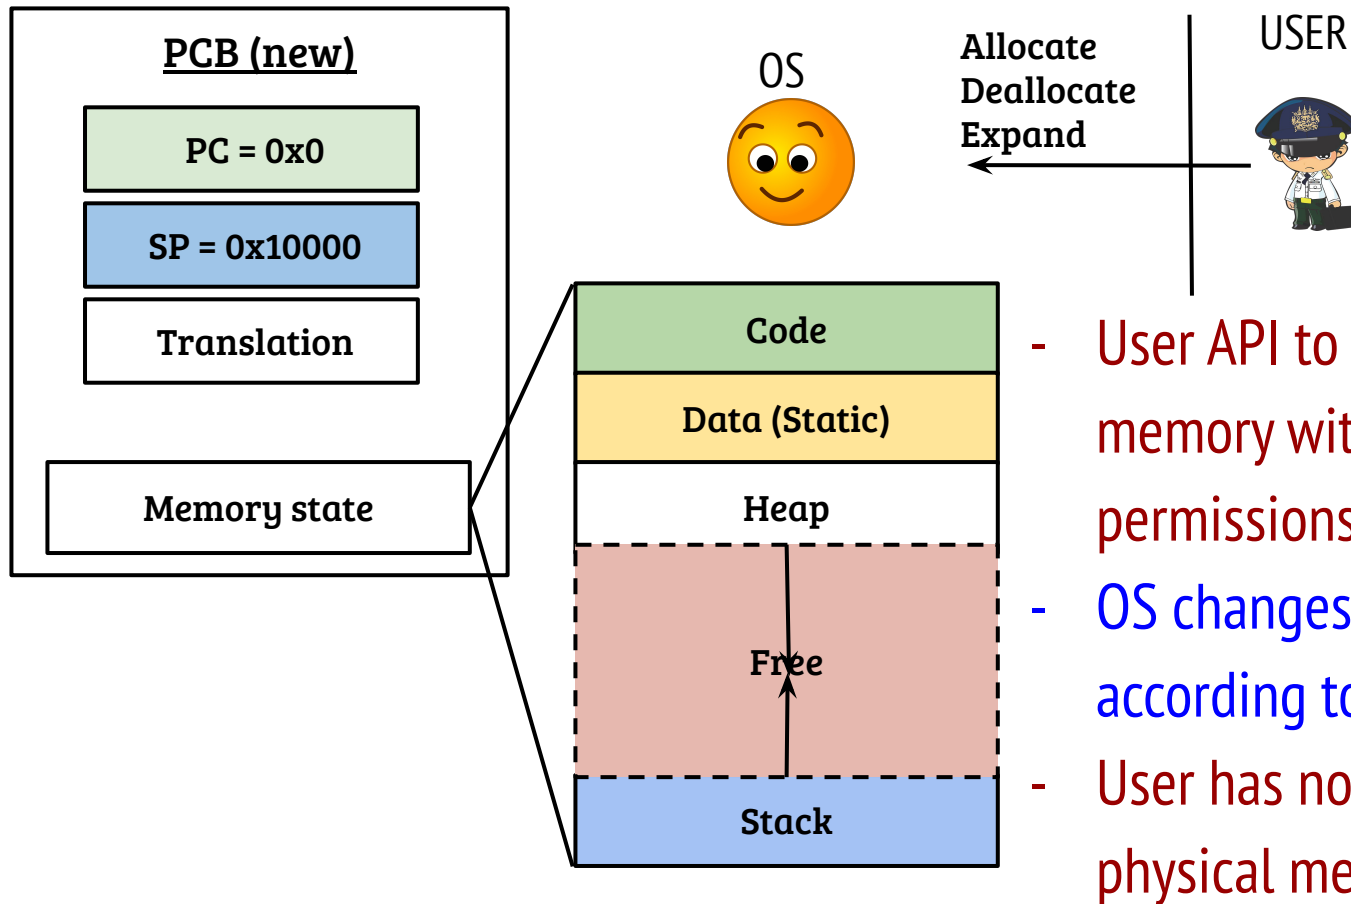


DRAM

The address space abstraction

- If all processes have same address space, how they map to actual memory?
- Architecture support used by OS techniques to perform memory virtualization i.e., translate virtual address to physical address (will revisit)
- What are the responsibilities of the OS during program load?
 - How CPU register state is changed?
- Creating address space, loading binary, updating the PCB register state
- What is the OS role in dynamic memory allocation?

User API for memory management

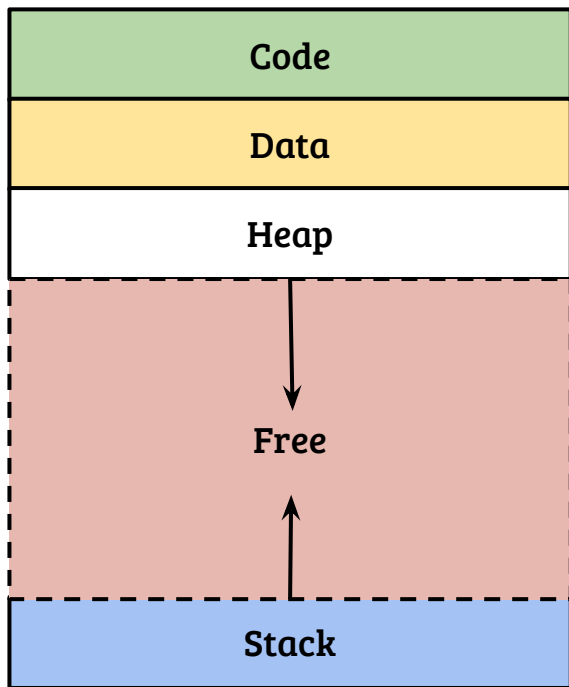


- User API to (de)allocate heap memory with different access permissions
- OS changes the memory state according to the user request
- User has no direct control on physical memory

CS330: Operating Systems

Virtual memory: Memory API

Recap: Process address space

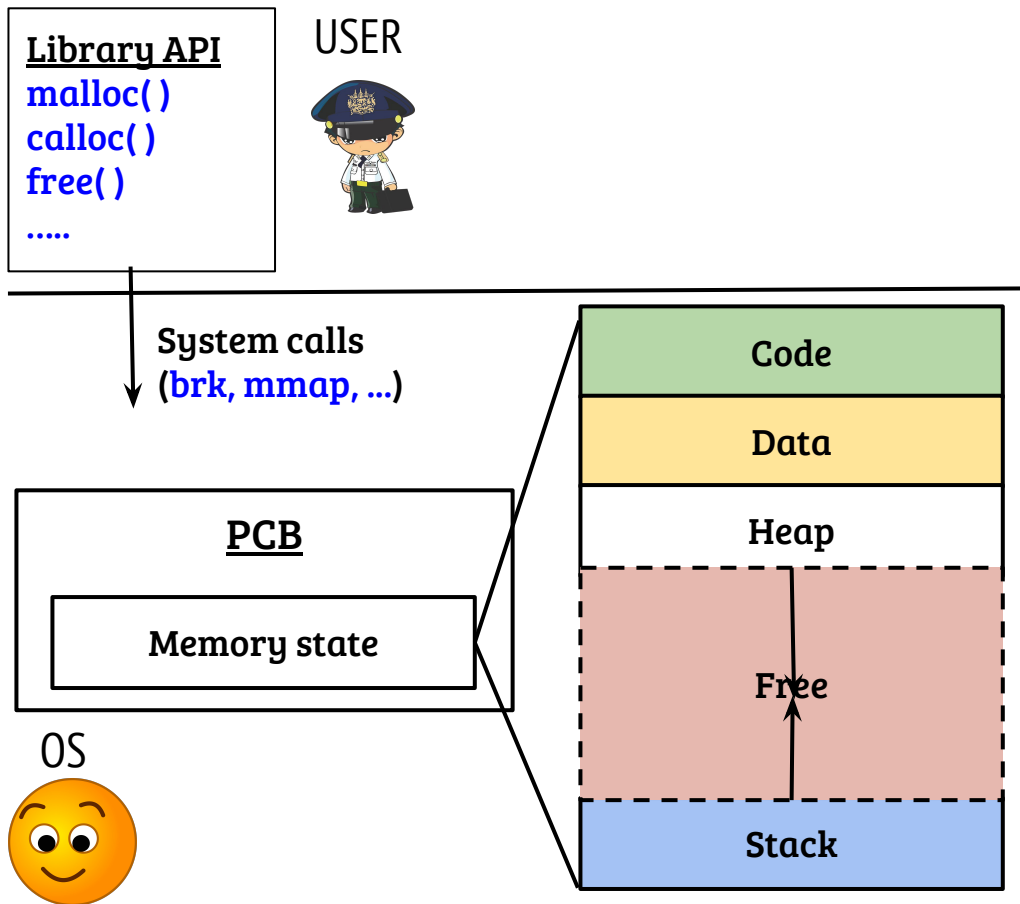


- Address space provides a unique view of memory to *all processes*
 - Address space is virtual
 - OS enables this virtual view

Recap: Process address space

- If all processes have same address space, how they map to actual memory?
- Architecture support used by OS techniques to perform memory virtualization i.e., translate virtual address to physical address (will revisit)
- What are the responsibilities of the OS during program load?
 - How CPU register state is changed?
- Creating address space, loading binary, updating the PCB register state
- What is the role of OS in dynamic memory allocation?

User API for memory management



- Generally, user programs use library routines to allocate/deallocate memory
- OS provides some address space manipulation system calls (today's agenda)

User API for memory management

Library API


`malloc()`
`calloc()`

USER




- Generally, user programs

- Can the size of segments change at runtime? If yes, which ones and how?
- How can we know about the segment layout at program load and runtime?
- How to allocate memory chunks with different permissions?
- What is the structure of PCB memory state?




calls (today's agenda)

Dynamically sizing the segments (UNIX)




- Code segment size and initialized data segment size is fixed (at exe load)

Dynamically sizing the segments (UNIX)



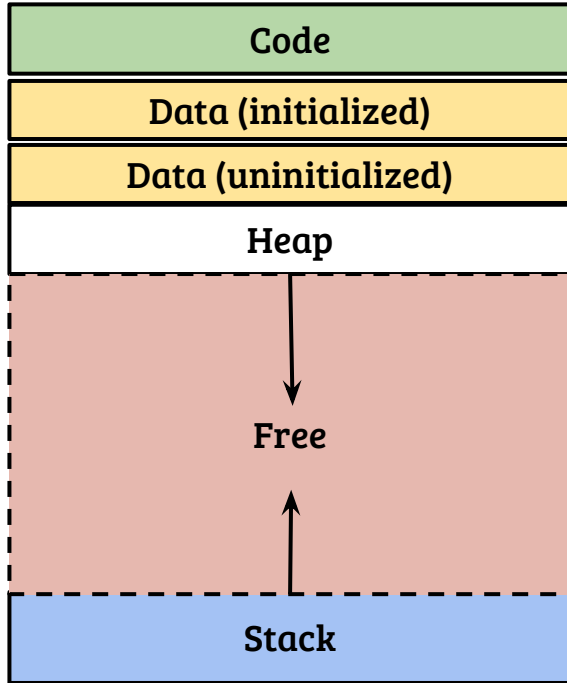
- Code segment size and initialized data segment size is fixed (at exe load)
- End of uninitialized data segment (a.k.a. BSS) can be adjusted dynamically

Dynamically sizing the segments (UNIX)



- Code segment size and initialized data segment size is fixed (at exe load)
- End of uninitialized data segment (a.k.a. BSS) can be adjusted dynamically
- Heap allocation can be discontinuous, special system calls like `mmap()` provide the facility

Dynamically sizing the segments (UNIX)



- Code segment size and initialized data segment size is fixed (at exe load)
- End of uninitialized data segment (a.k.a. BSS) can be adjusted dynamically
- Heap allocation can be discontinuous, special system calls like `mmap()` provide the facility
- Stack grows automatically based on the run-time requirements, no explicit system calls

Sliding the BSS (brk, sbrk)

```
int brk(void *address);
```

- If possible, set the end of uninitialized data segment (BSS) at address (address = heap start address)
- Can be used by C library to allocate/free memory dynamically

```
void * sbrk (long size);
```

- Increments the program's data space by size bytes and returns the old value of the end of bss
- `sbrk(0)` returns the current location of BSS

Finding the segments

- `etext`, `edata` and `end` variables mark the end of text segment, initialized data segment and the BSS, respectively (At program load)
- `sbrk(0)` can be used to find the start of heap segment
- Printing the address of functions and variables
- Linux provides the information in `/proc/pid/maps`

User API for memory management

Library API

USER

- Can the size of segments change at runtime? If yes, which ones and how?
- Heap and data segments can be adjusted using `brk` and `sbrk`
- How can we know about the segment layout at program load and runtime?
- Using predefined variables, `sbrk`, `proc` file system (Linux)
- How to allocate memory chunks with different permissions?
- What is the structure of PCB memory state?



Stack

Discontiguous allocation (mmap)

- `mmap()` is a powerful and multipurpose system call to perform dynamic and discontiguous allocation (explicit OS support)
- Allows to allocate address space
 - with different protections (READ/WRITE/EXECUTE)
 - at a particular address provided by the user
- Example: Allocate 4096 bytes with READ+WRITE permission

```
ptr = mmap(NULL, 4096, PROT_READ|PROT_WRITE, MAP_ANONYMOUS  
|MAP_PRIVATE, -1, 0); // See the man page for details
```

User API for memory management

Library API


USER

- Can the size of segments change at runtime? If yes, which ones and how?
- Heap and data segments can be adjusted using `brk` and `sbrk`
- How can we know about the segment layout at program load and runtime?
- Using predefined variables, `sbrk`, `proc` file system (Linux)
- How to allocate memory chunks with different permissions?
- `mmap()` supports discontinuous allocation with different permissions
- What is the structure of PCB memory state?



Stack

Memory state of PCB (example)



- Maintained as a sorted circular list accessible from PCB
- START and END never overlap between two segment areas
- Can merge/extend areas if permissions match

User API for memory management

Library API

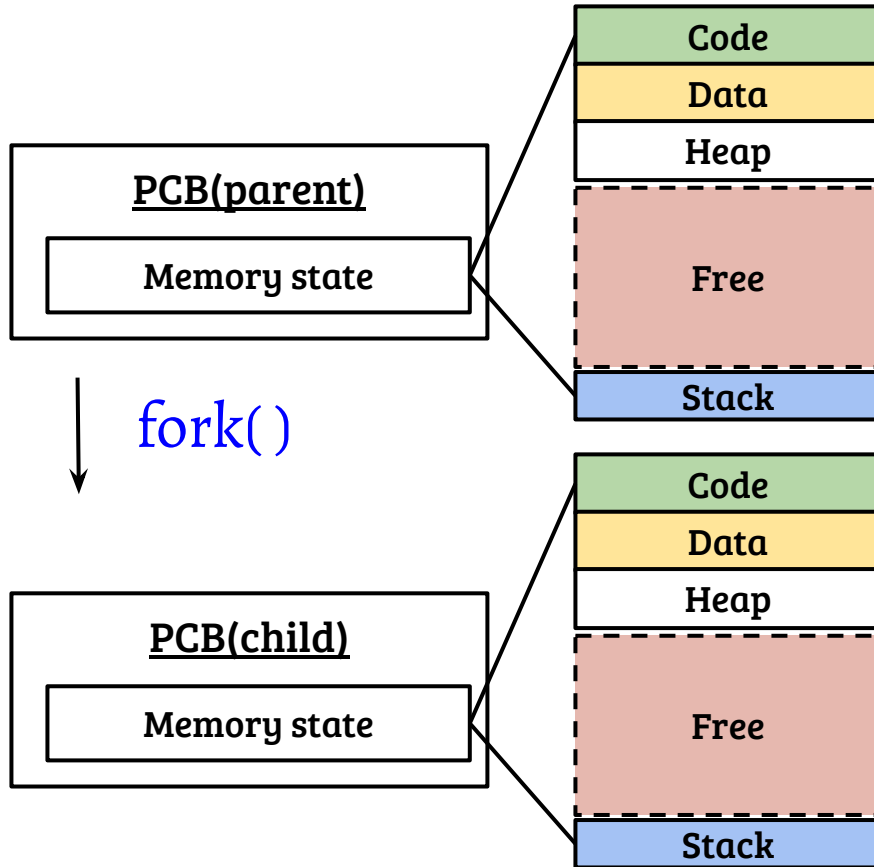
USER

- Can the size of segments change at runtime? If yes, which ones and how?
- Heap and data segments can be adjusted using `brk` and `sbrk`
- How can we know about the segment layout at program load and runtime?
- Using predefined variables, `sbrk`, `proc` file system (Linux)
- How to allocate memory chunks with different permissions?
- `mmap()` supports discontinuous allocation with different permissions
- What is the structure of PCB memory state?
- A sorted data structure of allocated areas can be used



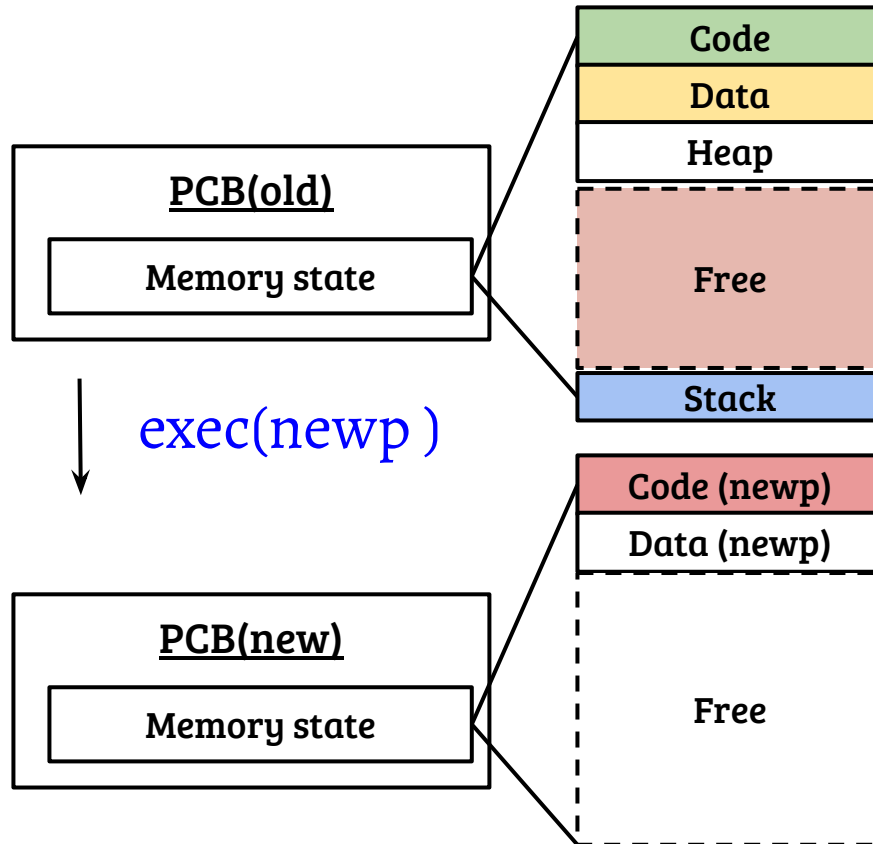
Stack

Inheriting address space through fork()



- Child inherits the memory state of the parent
 - The memory state data structures are copied into the child PCB
- Any change through `mmap()` or `brk()` is per-process

Overriding address space through exec()




- The address space is reinitialized using the new executable
- Changes to newly created address space depends on the logic of new process

CS330: Operating Systems


Virtual memory: Address translation

Recap: Process address space



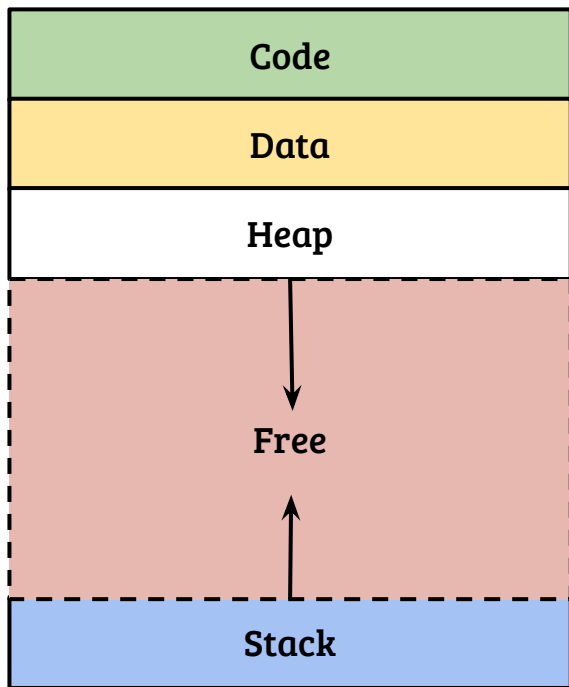
- Address space provides a unique view of memory to *all processes*
 - Address space is virtual
 - OS enables this virtual view

Recap: Process address space



- Address space provides a unique view of memory to *all processes*
 - Address space is virtual
 - OS enables this virtual view
- User can organize/manage virtual memory using OS APIs
 - No control on physical memory!


Recap: Process address space



- Address space provides a unique view of memory to *all processes*
 - Address space is virtual
 - OS enables this virtual view
- User can organize/manage virtual memory using OS APIs
 - No control on physical memory!

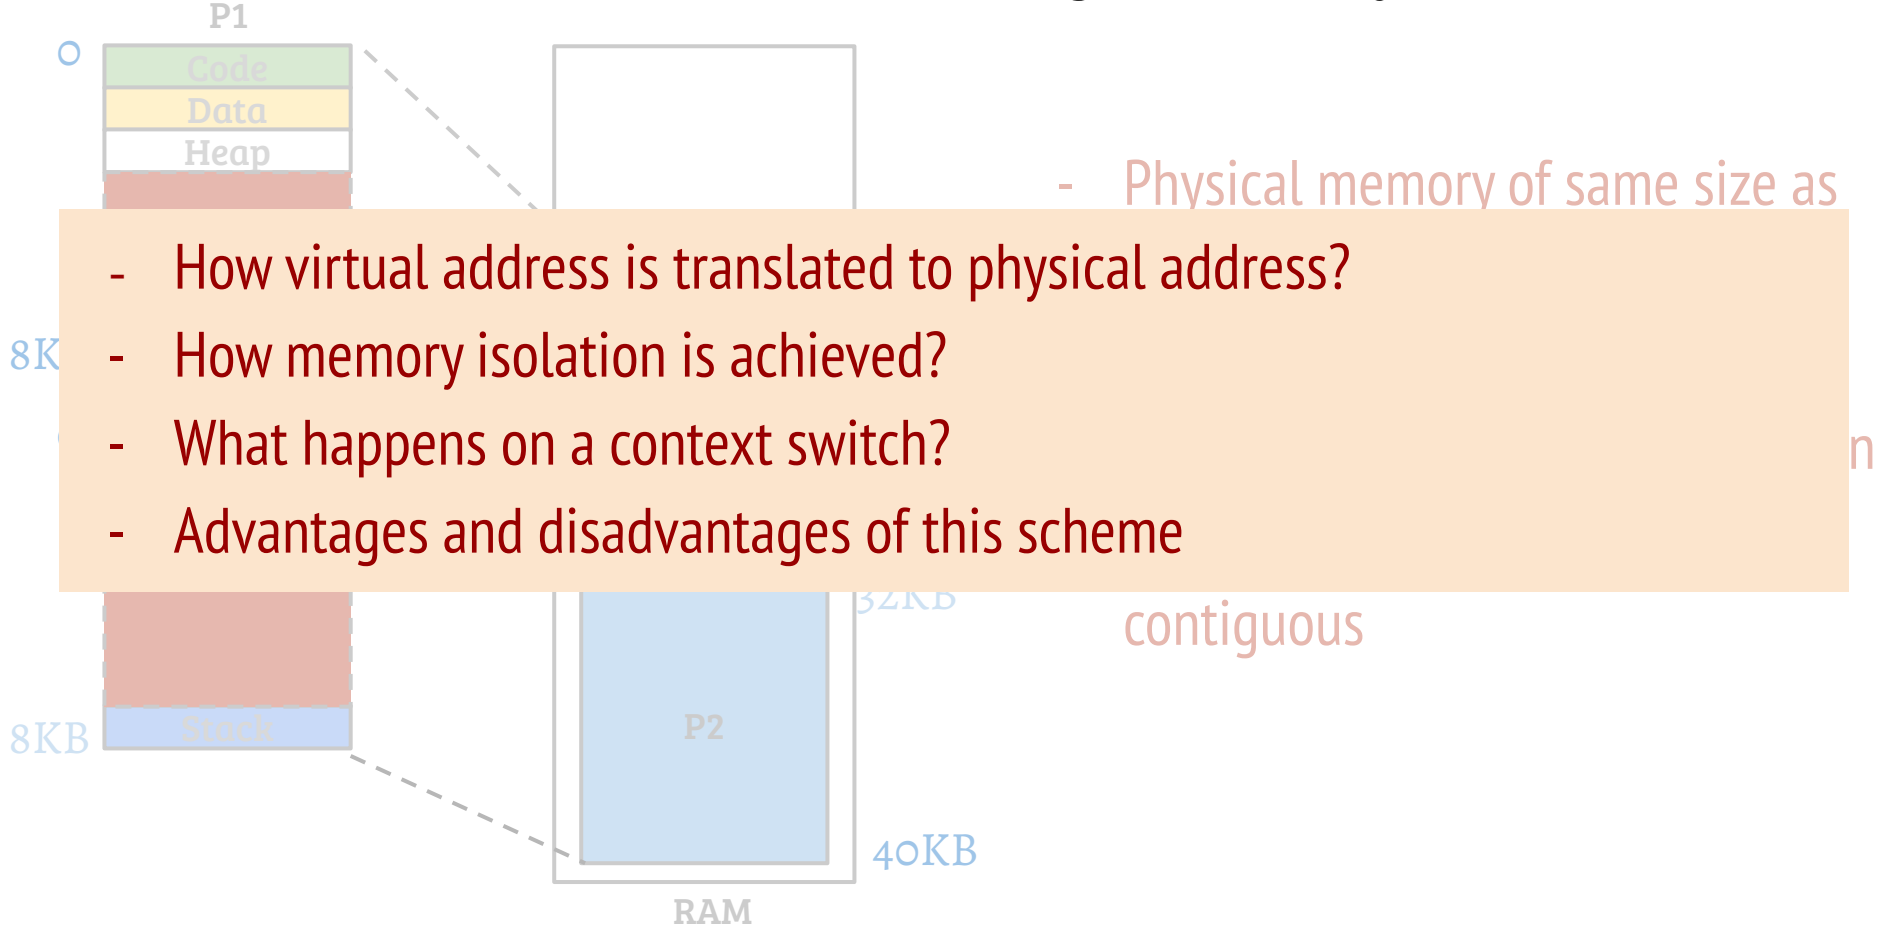
Today's agenda: Virtual to physical address translation

Translation at address space granularity



- Physical memory of same size as the address space is allocated to each process
- Physical memory for a process can be at any address, but should be contiguous

Translation at address space granularity



Role of the compiler

Simple function

```
func()  
{  
    int a = 100;  
    a+ = 10;  
}
```

Compiled assembly

.....

func:

```
10:  push %rbp;  
12:  mov %rsp, %rbp;  
15:  mov (%rbp), %rax  
18:  add $10, %rax  
21:  mov %rax, (%rbp)  
24:  pop %rbp;  
26:  ret;
```

.....

- Compiler generates the code with starting address zero
- Compiler does not know the stack address, blindly uses the registers (rbp, rsp)!

OS during binary load (simplified exec)

```
load_new_executable( PCB *current, File *exe)
```

```
{
```

```
    reinit_address_space(current → mm_state);
```

```
    allocate_phys_mem(current);
```


```
    load_exe_to_physmem(current, exe);
```

```
    set_user_sp(current → mm_state → stack_start);
```

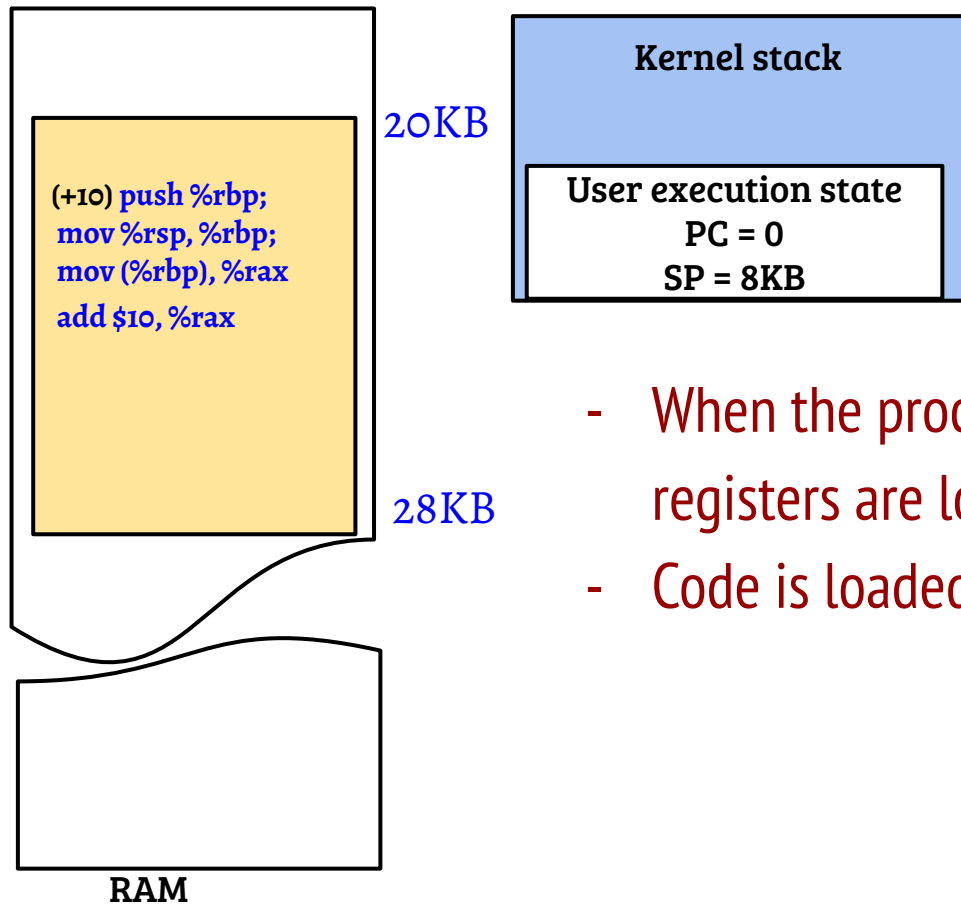
```
    set_user_pc(current → mm_state → code_start);
```

```
    return_to_user;
```

```
}
```

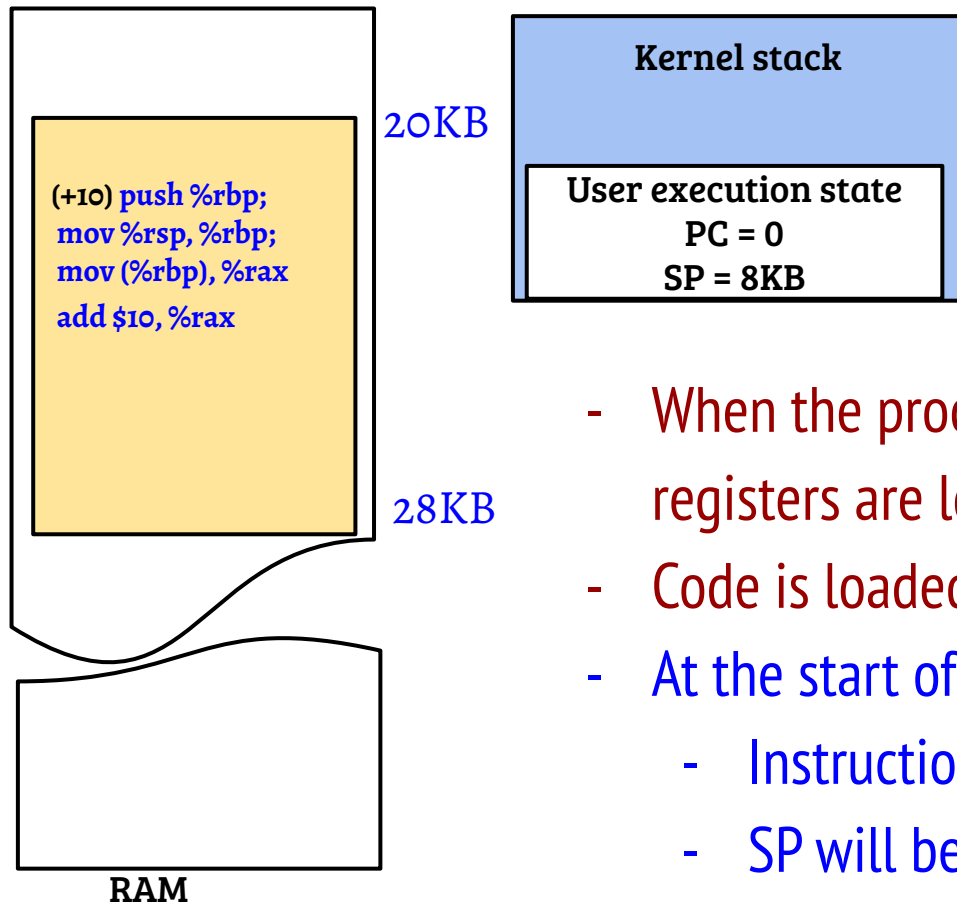


Process state after exec()





- When the process returns to user space, the registers are loaded with virtual addresses
- Code is loaded into physical memory (@20KB)

Process state after exec()



- When the process returns to user space, the registers are loaded with virtual addresses
- Code is loaded into physical memory (@20KB)
- At the start of “func” execution
 - Instruction fetch address is 10 (PC = 10)
 - SP will be around 8KB

Process state after exec()






Dear HW! I have done my part. Help me with the translation, please!






- When the process returns to user space, the registers are loaded with virtual addresses
- Code is loaded into physical memory (@20KB)
- At the start of "func" execution
 - Instruction fetch address is 10
 - RSP will be around 8KB

Process state after exec()



Here is a base register. I will add the value of base register with the virtual address generated by the program to get the physical address. All yours buddy!

Process state after exec()





Here is a base register. I will add the value of base register with the virtual address generated by the program to get the physical address. All yours buddy!





Hurray! I will configure the value of base register as per my need. I see some light atlast!

Translation




- In this case, base register value should be 20KB
- InsFetch (vaddr = 10) \Rightarrow InsFetch (paddr = 20KB + 10)
- How "push %rbp" works?

Translation

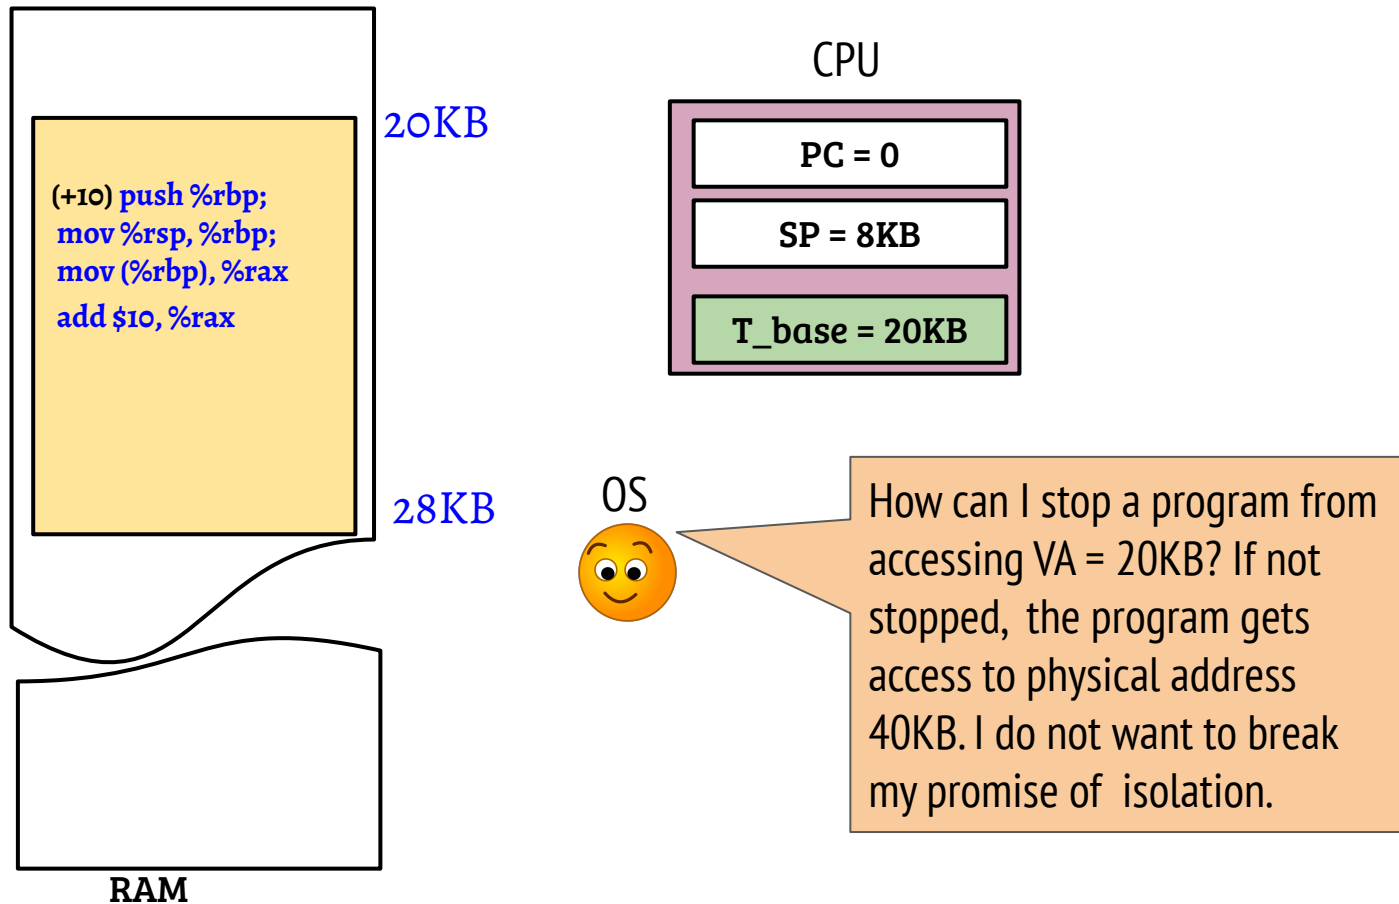


- In this case, base register value should be 20KB
- InsFetch (vaddr = 10) \Rightarrow InsFetch (paddr = 20KB + 10)
- How “push %rbp” works?
- Assuming RSP = 8KB, “push %rbp” results in a memory store at address (8KB - 8)
 - CPU translates the address to (28KB - 8)



Translation at address space granularity



Isolation: How to stop illegal access?


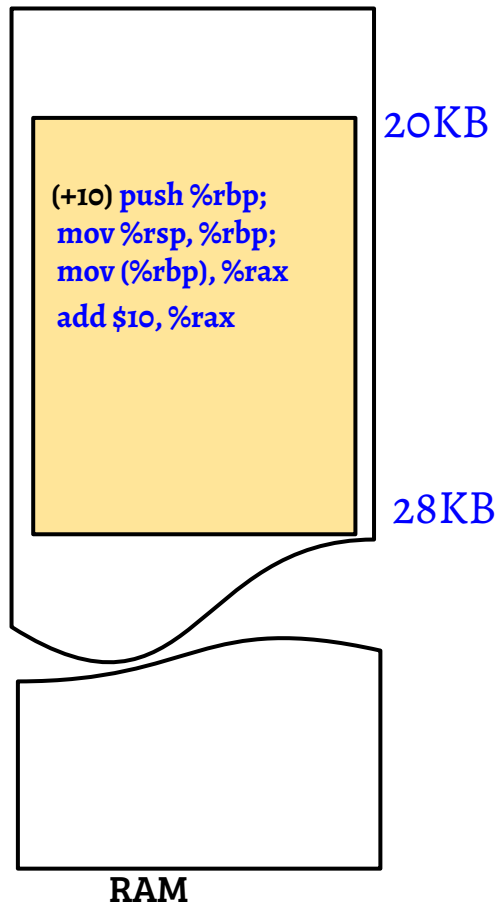


Isolation: How to stop illegal accesses?



Once a cry baby always a cry baby! I also provide a limit register to enforce the limit during translation. Before you ask, these registers can only be changed from privileged mode

Isolation: How to stop illegal accesses?



Once a cry baby always a cry baby! I also provide a limit register to enforce the limit during translation. Before you ask, these registers can only be changed from privileged mode

- The hardware raises a fault if some program violates the limit.
- The OS fault handler may kill the process

Translation at address space granularity



- How virtual address is translated to physical address?
- The OS sets the base register value depending on the physical location.
The hardware performs the translation using the base value.
- How memory isolation is achieved?
- Limit register can be used to enforce memory isolation
- What happens on a context switch?
- Advantages and disadvantages of this scheme



Context switch and translation information

- The base and limit register values can be saved in the outgoing process PCB during context switch
- Loaded from PCB to the CPU when a process is scheduled

Translation at address space granularity

P1

- How virtual address is translated to physical address?
- The OS sets the base register value depending on the physical location.
The hardware performs the translation using the base value.
- How memory isolation is achieved?
- Limit register can be used to enforce memory isolation
- What happens on a context switch?
- Save and restore limit and base registers
- Advantages and disadvantages of this scheme

RAM

Translation at address space granularity: Issues

- Physical memory must be greater than address space size
 - Unrealistic, against the philosophy of address space abstraction
 - Small address space size \Rightarrow Unhappy user


Translation at address space granularity: Issues

- Physical memory must be greater than address space size
 - Unrealistic, against the philosophy of address space abstraction
 - Small address space size \Rightarrow Unhappy user
- Memory inefficient
 - Physical memory size is same as address space size irrespective of actual usage \Rightarrow Memory wastage
 - Degree of multiprogramming is very less

CS330: Operating Systems


Virtual memory: Segmentation

Recap: Translation at address space granularity



- Physical memory of same size as the address space size is allocated to each process
- Issues: Memory inefficient, inflexible


Recap: Translation at address space granularity





- Physical memory of same size as the address space size is allocated to each process
- Issues: Memory inefficient, inflexible

Today's agenda: Translation at segment granularity

Segmentation




Address space



RAM

- Extension of the basic scheme with more base-limit register pairs

Segmentation



Address space

Code segment

Base = 28KB

Limit = 1KB

Data segment


Base = 20KB

Limit = 2KB

Stack segment

Base = 32KB

Limit = 1KB




RAM

- Example
 - Code address
 - Data address

Segmentation



- How the CPU decides which segment to use?
- How stack growth in opposite direction handled?
- What happens on context switch?
- Advantages and disadvantages of segmentation



Segmentation: Explicit addressing

- Part of the address is used to explicitly specify segments
- In our example,
 - virtual address space = 8KB, address length = 13 bits and there are three segments
 - Two MSB bits used to specify the segment: “00” for code, “01” for data and “11” for stack
 - The hardware selects the segment register based on the value of two MSB bits and rest of the bits are used as the offset
 - Max. size of each segment = 2KB

Issues with explicit addressing

- Inflexible
 - Data and stack can not be sized dynamically
- Wastage of virtual address space
 - In our example, 2KB virtual address is unusable
- Note: Physical allocation is still done in an on-demand basis


Segmentation: Implicit addressing

- The hardware selects the segment register based on the operation
- Code segment for instruction access
 - Fetch address, jump target, call address


Segmentation: Implicit addressing

- The hardware selects the segment register based on the operation
- Code segment for instruction access
 - Fetch address, jump target, call address
- Stack segment for stack operations
 - Arguments for push and pop, indirect addressing with SP, BP
- Data segment for other addresses

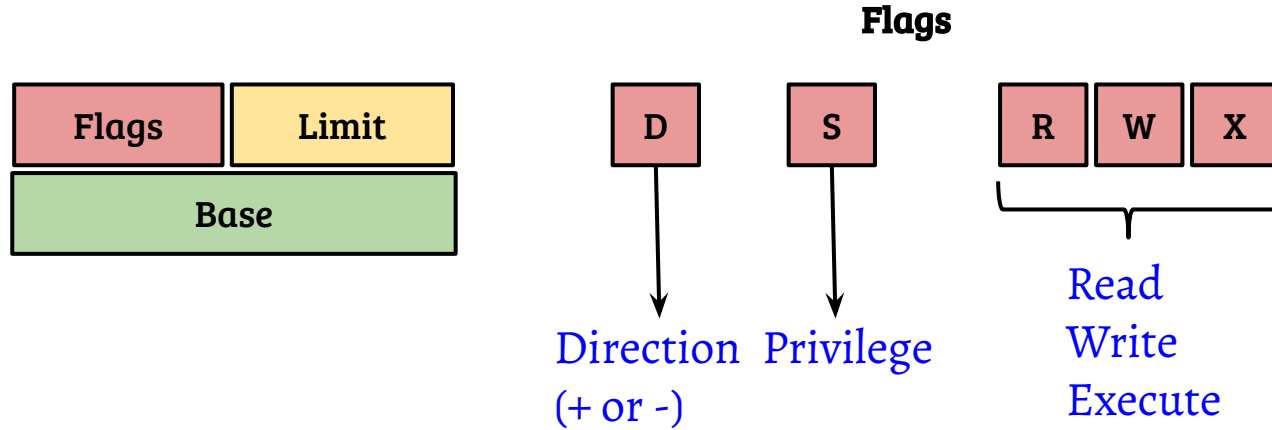
Segmentation



- How the CPU decides which segment to use?
- Explicit and implicit addressing
- How stack growth in opposite direction handled?
- What happens on context switch?
- Advantages and disadvantages of segmentation




Segmentation (protection and direction)



- For stack, direction is -ve, used by hardware to calculate physical address
- “S” bit can be used to specify privilege, specifically useful in code segment
- R, W and X can be used to enforce isolation and sharing

Segmentation

- How the CPU decides which segment to use?
- Explicit and implicit addressing
- How stack growth in opposite direction handled?
- Flag bits for direction of growth, access permissions
- What happens on context switch?
- Save and restore segment registers
- Advantages and disadvantages of segmentation



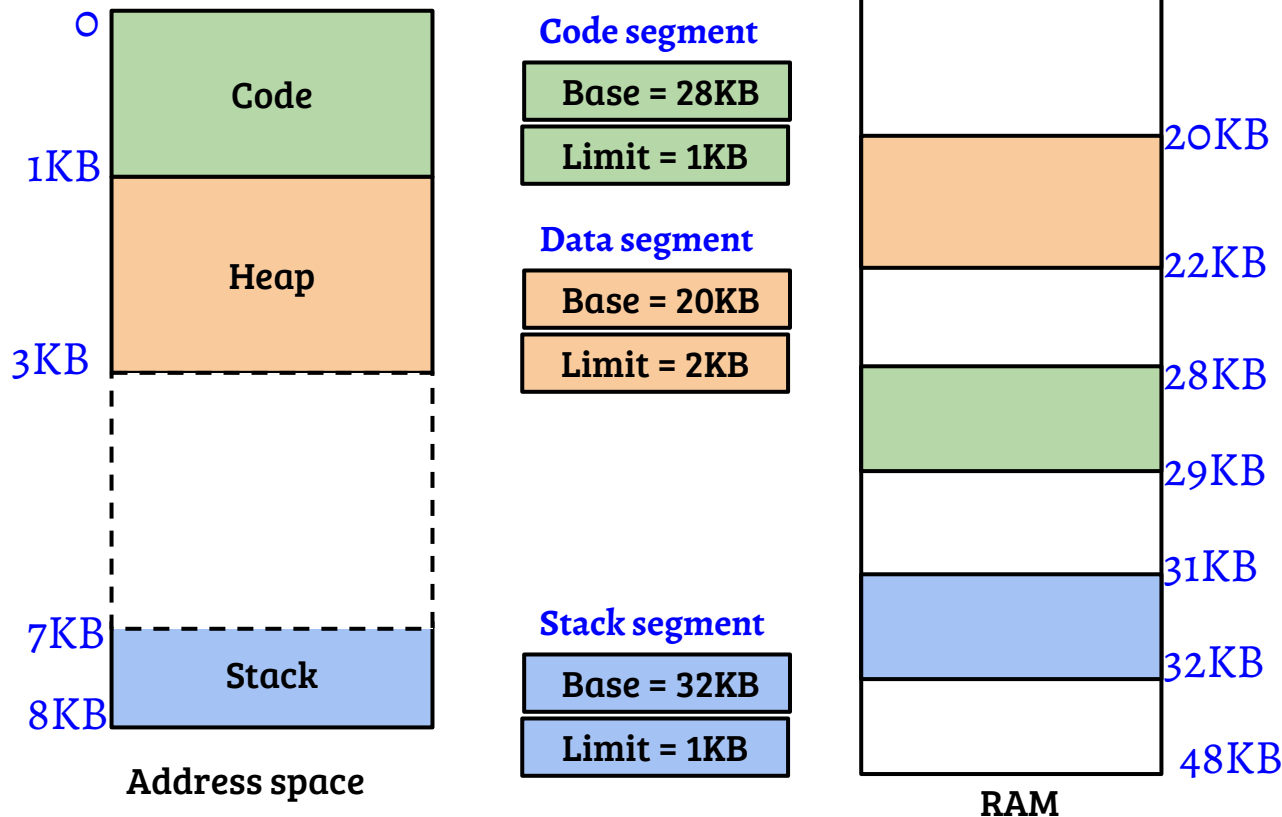
Advantages and disadvantages of segmentation

- Advantages
 - Easy and efficient address translation
 - Save memory wastage for unused addresses
- Disadvantages
 - External fragmentation
 - Can not support discontinuous sparse mapping

CS330: Operating Systems


Virtual memory: Paging

Recap: Segmentation



- Extension of the scheme for translation and address space granularity
- Base-limit register pairs per segment

Recap: Segmentation in reality



- Descriptor table register (DTR) is used to access the descriptor table
- # of descriptors depends on architecture
- Separate descriptors used for user and kernel mode

Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it page)
 - Physical memory partitioned in a similar way (call it page frame)


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it pages)
 - Physical memory partitioned in a similar way (call it page frames)
 - OS creates a mapping between *page* to *page frame*
 - H/W uses the mapping to translate VA to PA


Paging example (pages)




- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128




Paging example (pages)



- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128




Virtual address





- Example: For Virtual address *0x0510*, Page number = 5, offset = 16


Paging example (page frames)



- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



DRAM



Process address
space

32KB


Paging example (page frames)

0	Page 0
256	Page 1
512	Page 2
768	Page 3
1024	

Page 125
Page 126
Page 127

Process address
space

- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



Physical address

PFN 0	0
PFN 1	256
PFN 2	512
PFN 3	768
	1024


PFN 253
PFN 254
PFN 255

64KB

DRAM



- Example: For physical address *0x1F51*, PFN = 31, offset = 81

Paging example (page table mapping)





- Each entry in page table is called page table entry (PTE)
- Example mapping: page 0 \Rightarrow PFN 1, page 2 \Rightarrow PFN 2 and so on


Paging example (page table walk)



```
PTW (vaddr V, PTable P)
// Input: Virtual address, Page table
// Returns physical address
{
    Entry = P[V >> 8];
    if (Entry.present)
        return (Entry.PFN << 8) + (V & 0xFF);
    Raise PageFault;
}
```




Paging example (example translation)

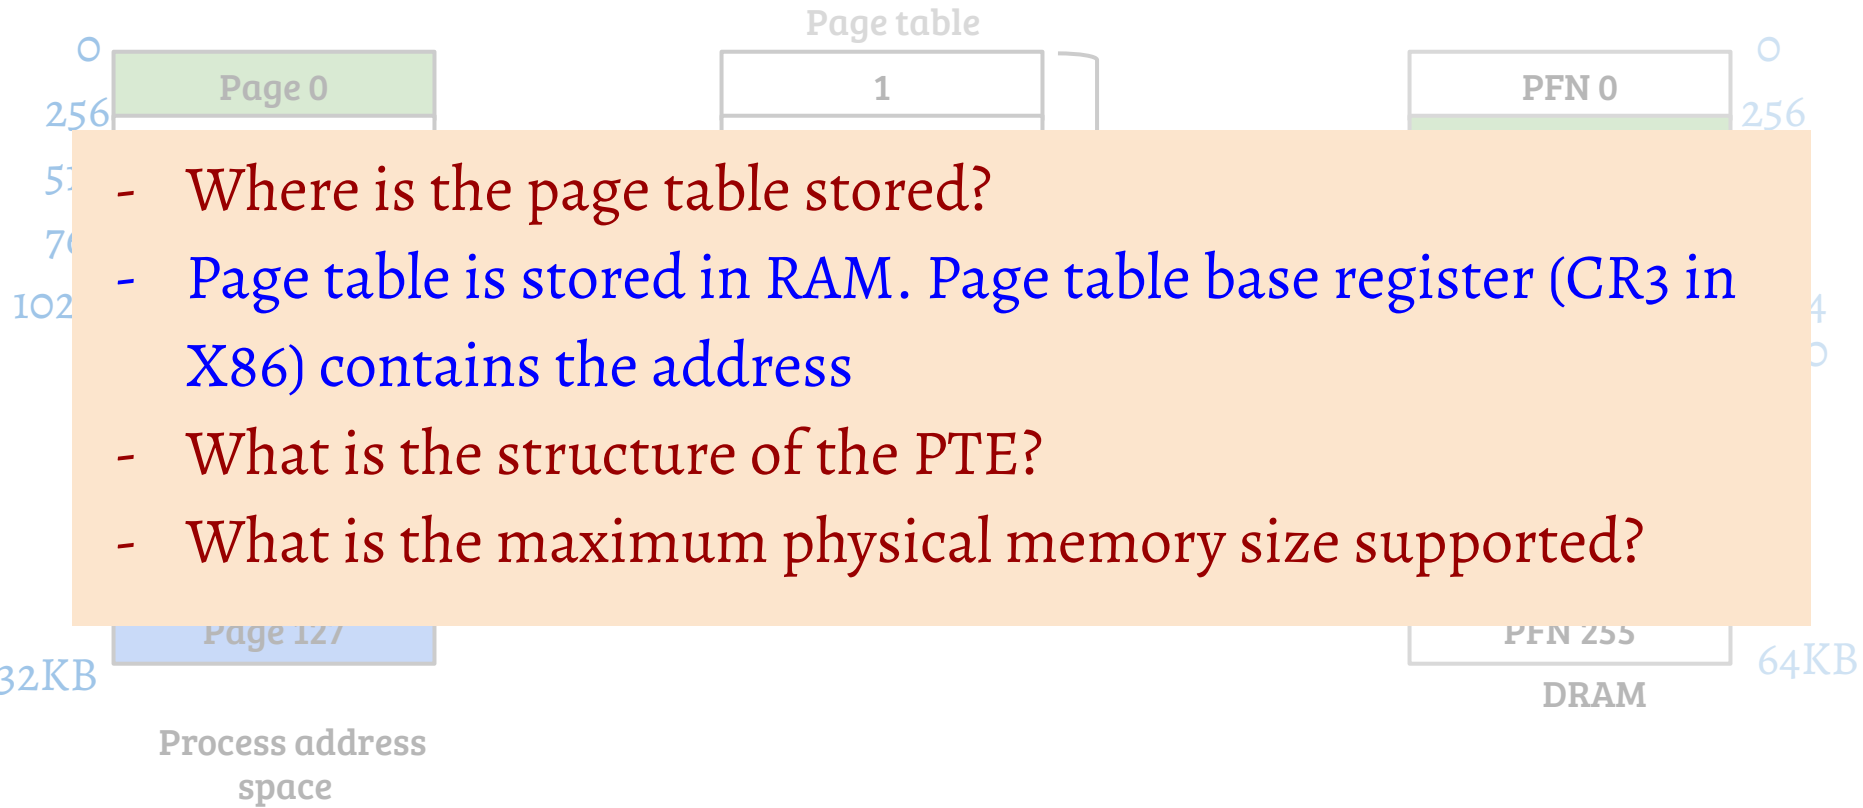


- Process address space
- Virtual address 0x10 translates to physical address 0x110
 - Virtual address 0x7FF0 translates to physical address 0x3F0

Paging example (page table walk)



Paging example (page table walk)



Paging example (structure of an example PTE)



- PFN occupies a significant portion of PTE entry (8 bits in this example)

P

Present bit, 1 \Rightarrow entry is valid

W

Write bit, 1 \Rightarrow Write allowed

S

Privilege bit, 0 \Rightarrow only kernel mode access is allowed

A

Accessed bit, 1 \Rightarrow Address accessed (set by H/W during walk)

D

Dirty bit, 1 \Rightarrow Address written (set by H/W during walk)

X

Execute bit, 1 \Rightarrow Instruction fetch allowed for this page

Reserved/unused bits

Paging example (Page table entries)

0	Page 0
256	Page 1
512	Page 2
768	Page 3
1024	

Page table	
0x125	
0x0	
0x207	
0x407	
0x0	

0	PFN 0	
256	PFN 1	
512	PFN 2	
768	PFN 3	
1024	PFN 4	
1280		

Page 125
Page 126
Page 127

0x0
0x307

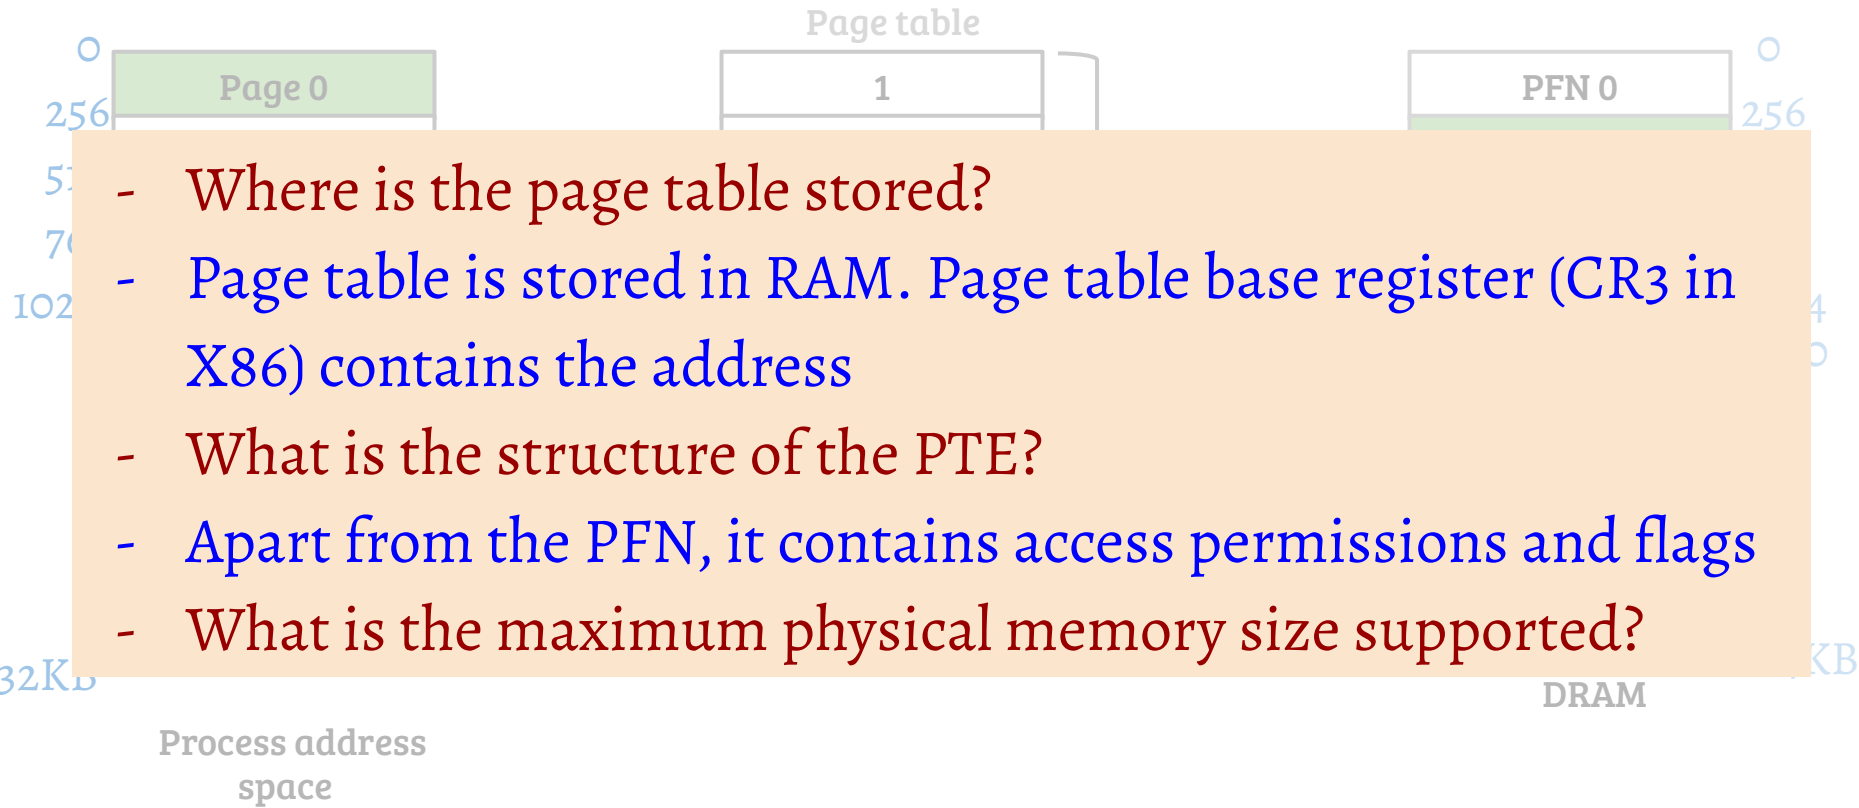
PFN 253
PFN 254
PFN 255

64KB

Process address
space

- Code: Page 0 (Read and Execute)
- Data: Page 2 and Page 3 (Read and Write)
- Stack: Page 127 (Read and Write)

Paging example (page table walk)



Paging example (page table walk)

- Where is the page table stored?
- Page table is stored in RAM. Page table base register (CR3 in X86) contains the address
- What is the structure of the PTE?
- Apart from the PFN, it contains access permissions and flags
- What is the maximum physical memory size supported?
- For this example, 8-bits can be used to specify 256 page frames. Maximum RAM size = $256 * 256 = 64\text{KB}$

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?

Paging: one level of page table may not be feasible!

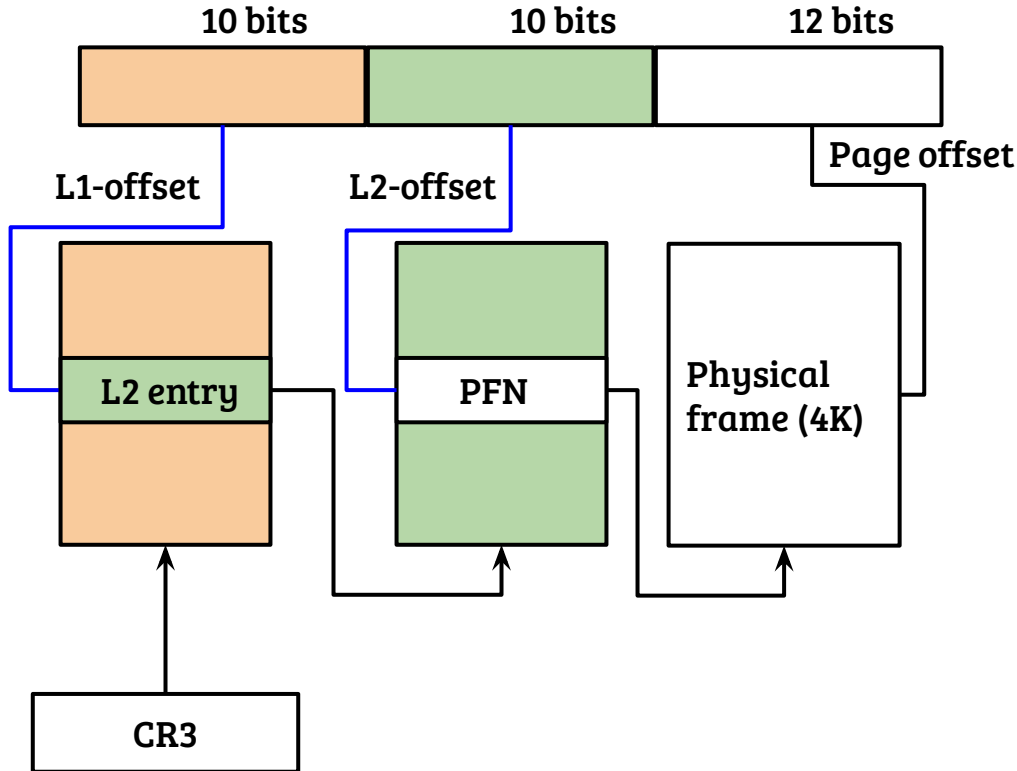
- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system?

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system? (2^{20} entries)
- Not possible to hold 2^{20} entries in a single page
- Therefore, multi-level page tables are used in modern systems

Two-level page tables (32-bit virtual address)

Virtual Address



- Two-level page table
- Level-1 page table contains entries pointing to Level-2 page table structures
- Level-2 entry contains PFN along with flags

CS330: Operating Systems


Virtual memory: Multilevel paging and TLB

Recap: Paging

- The idea of paging
 - Partition the address space into fixed sized blocks (call it pages)
 - Physical memory partitioned in a similar way (call it page frames)
 - OS creates a mapping between *page* to *page frame* , H/W uses the mapping to translate VA to PA
- With increased address space size, single level page table entry is not feasible, because
 - Increasing page size increases internal fragmentation
 - Small pages may not be suitable to hold all mapping entries


Today's agenda: Multi-level pages tables and implications

4-level page tables: 48-bit VA (Intel x86_64)



- Virtual address size = 2^{48} , Page size = 4096 bytes
- Four-levels of page table, entry size = 64 bits

4-level page tables: example translation



- Hardware translation by repeated access of page table stored in physical memory
- Page table entry: 12 bits LSB is used for access flags

Paging: translation efficiency

	0x20100: mov \$0, %rax;	
	0x20102: mov %rax, (%rbp);	// sum=0
sum = 0;	0x20104: mov \$0, %rcx;	// ctr=0
for(ctr=0; ctr<10; ++ctr)	0x20106: cmp \$10, %rcx;	// ctr < 10
sum += ctr;	0x20109: jge 0x2011f;	// jump if >=
	0x2010f: add %rcx, %rax;	
	0x20111: mov %rax, (%rbp);	// sum += ctr
	0x20113: inc %rcx	// ++ctr
	0x20115: jmp 0x20106	// loop
	0x2011f:	

- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

Paging: translation efficiency

```
0x20100: mov $0, %rax;
```

```
0x20102: mov %rax, (%rbp); // sum=0
```

- Instruction execution: Loop = $10 * 6$, Others = $2 + 3$
 - Memory accesses during translation = $65 * 4 = 260$
 - Data/stack access: Initialization = 1, Loop = 10
 - Memory accesses during translation = $11 * 4 = 44$
 - A lot of memory accesses (> 300) for address translation
 - How many distinct pages are translated?
-
- Considering four-level page table, how many memory accesses are required (for translation) during the execution of the above code?

Paging with TLB: translation efficiency

Translate(V){

PageAddress P = V >> 12;

TLBEntry entry = lookup(P);

if (entry.valid) return entry.pte;

entry = PageTableWalk(V);

MakeEntry(entry);

return entry.pte;

}

TLB	
Page	PTE
0x20	0x750
0x7FFF	0x890

- TLB is a hardware cache which stores *Page* to *PFN* mapping
- After first miss for instruction fetch address, all others result in a TLB hit
- Similarly, considering the stack virtual address range as 0x7FFF000 - 0x8000000, one entry in TLB avoids page table walk after first miss


Paging: translation efficiency

```
0x20100: mov $0, %rax;
```

- Instruction execution: Loop = $10 * 6$, Others = $2 + 3$
 - Memory accesses during translation = $65 * 4 = 260$
- Data/stack access: Initialization = 1, Loop = 10
 - Memory accesses during translation = $11 * 4 = 44$
- A lot of memory accesses (> 300) for address translation
- How many distinct pages are translated?
- One code page (0x20) and one stack page (0x7FFF). Caching these translations, will save a lot of memory accesses.


required (for translation) during the execution of the above code?

Address translation (TLB + PTW)




- TLB in the path of address translation
- Separate TLBs for instruction and data, multi-level TLBs
- In X86, OS can not make entries into the TLB directly, it can flush entries

Address translation (TLB + PTW)



- TLB in the path of address

- How TLB is shared across multiple processes?
- Why page fault is necessary?
- How OS handles the page fault?



into the TLB directly, it can flush entries

TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
 - A) Do nothing
 - B) Flush the whole TLB
 - C) Some other solution

TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
 - A) Do nothing
 - B) Flush the whole TLB
 - C) Some other solution
- Process B may be using the same addresses used by A. Result: Wrong translation

TLB: Sharing across applications

Process (A)

Process (B)

Page	PTE
0x100	0x200007
0x101	0x205007

TLB

- Assume that, process A is currently executing. What happens when process B is scheduled?
 - A) Do nothing
 - B) Flush the whole TLB
 - C) Some other solution
- Correctness ensured. Performance is an issue (with frequent context switching)

TLB: Sharing across applications

Process (A)

Process (B)

- Assume that, process A is currently executing. What happens when process B is scheduled?


- A) Do nothing
- B) Flush the whole TLB
- C) Some other solution

- Address space identified (ASID) along with each TLB entry to identify the process

ASID	Page	PTE
A	0x100	0x200007
A	0x101	0x205007
B	0x100	0x301007
B	0x101	0x302007


TLB

Address translation (TLB + PTW)



- TLB in the path of address

- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- How OS handles the page fault?




entries

Address translation (TLB + PTW)



- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?



Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Error code

Other and unused	I	R	U	W	P
------------------	---	---	---	---	---

P

Present bit, 1 \Rightarrow fault is due to protection

W

Write bit, 1 \Rightarrow Access is write

U

Privilege bit, 1 \Rightarrow Access is from user mode

R

Reserved bit, 1 \Rightarrow Reserved bit violation

I

Fetch bit, 1 \Rightarrow Access is Instruction Fetch

- Error code is pushed into the kernel stack by the hardware

Page fault handling in X86: OS fault handler

```
HandlePageFault( u64 address, u64 error_code)
{
    If( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn( );
        install_pte(address, PFN);
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

Address translation (TLB + PTW)

- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?
- The hardware invokes the page fault handler by placing the error code and virtual address. The OS handles the page fault either fixing it or raising a SEGFault.


CS330: Operating Systems

Virtual memory: Page fault and Swapping

Recap: Address translation



- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?



Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Page fault handling in X86: Hardware

```
If( !pte.valid ||  
    (access == write && !pte.write) ||  
    (cpl != 0 && pte.priv == 0)){  
    CR2 = Address;  
    errorCode = pte.valid  
                | access << 1  
                | cpl << 2;  
    Raise pageFault;  
} // Simplified
```

Error code

Other and unused	I	R	U	W	P
------------------	---	---	---	---	---

P

Present bit, 1 \Rightarrow fault is due to protection

W

Write bit, 1 \Rightarrow Access is write

U

Privilege bit, 1 \Rightarrow Access is from user mode

R

Reserved bit, 1 \Rightarrow Reserved bit violation

I

Fetch bit, 1 \Rightarrow Access is Instruction Fetch

- Error code is pushed into the kernel stack by the hardware

Page fault handling in X86: OS fault handler


```
HandlePageFault( u64 address, u64 error_code)
{
    If( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn( );
        install_pte(address, PFN);
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

Address translation (TLB + PTW)

- How TLB is shared across multiple processes?
- Full TLB flush during context switch, using ASID
- Why page fault is necessary?
- Page fault is required to support memory over-commitment through lazy allocation and swapping
- How OS handles the page fault?
- The hardware invokes the page fault handler by placing the error code and virtual address. The OS handles the page fault either fixing it or raising a SEGFAULT.

Swapping (swap-out)

DRAM



Swap (Hard disk)


Number of free PFNs are very few in the system. I can not break my promise made to the applications. Let me swap-out some memory. But which one to swap-out?



OS

AllocatePFN()

Swapping (swap-out)



My page replacement policy will help me deciding the victims (V). Can I just swap-out? What if the swapped-out pages are accessed? I should be prepared for that too!



OS

AllocatePFN()

Swapping (swap-out)


DRAM



PTE mapping the victim PFN (before swap)



PTE mapping the victim PFN (after swap)



Swap (Hard disk)

Update the present-bit to 0 in the PTE such that any access to the page through the virtual address will result in a page fault. Also maintain the swap address in the PTE.



OS

AllocatePFN()

Swapping (swap-out)


DRAM



PTE mapping the victim PFN (before swap)



PTE mapping the victim PFN (after swap)



Swap (Hard disk)

Content of the PFN is now in the swap device. In future, any translation using the PTE will result in a page fault. The page fault handler would copy it back from the swap device.



OS

AllocatePFN()

Page fault: Swap-in

```
HandlePageFault( u64 address, u64 error_code)
{
    If ( AddressExists(current → mm_state, address) &&
        AccessPermitted(current → mm_state, error_code) {
        PFN = allocate_pfn();
        If ( is_swapped_pte(address) )    // Check if the PTE is swapped out
            swapin(getPTE(address), PFN); // Copy the swap block to PFN
        install_pte(address, PFN);       // and update the PTE
        return;
    }
    RaiseSignal(SIGSEGV);
}
```

Page replacement

- Objective: minimize number of page faults (due to swapping)
- We can model this problem with three parameters
 - A given sequence of access to virtual pages
 - # of memory pages (Frames)
 - Page replacement policy
- Metrics to measure the effectiveness: # of page faults, page fault rate, average memory access time

Belady's optimal algorithm (MIN)

- Strategy: Replace the page that will be referenced after the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = ?

Belady's optimal algorithm (MIN)

- Strategy: Replace the page that will be referenced after the longest time
- Example:
 - #of frames = 3
 - Reference sequence (in temporal order)
1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3
- #of page faults = 6 (3 cold-start misses result in page faults, no swapping)
- Belady's MIN is proven to be optimal, but impractical as it requires knowledge of future access

First In First Out (FIFO)

- Strategy: Replace the page that is in memory for the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = ?

First In First Out (FIFO)

- Strategy: Replace the page that is in memory for the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = 8 (3 cold-start misses)
- FIFO suffers from an anomaly known as Belady's anomaly
 - With increased #of frames, #of page fault may also increase!

First In First Out (FIFO)

- Strategy: Replace the page that is in memory for the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = 8 (3 cold-start misses)
- FIFO suffers from an anomaly known as Belady's anomaly
 - With increased #of frames, #of page fault may also increase!
 - Example access sequence: 0, 1, 2, 3, 0, 1, 4, 0, 1, 2, 3, 4
 - #of page faults with 3 frames < #of page faults with 4 frames

Least recently used (LRU)

- Strategy: Replace the page that is not referenced for the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = ?

Least recently used (LRU)

- Strategy: Replace the page that is not referenced for the longest time

- Example:

#of frames = 3

Reference sequence (in temporal order)

1, 3, 1, 5, 4, 1, 2, 5, 2, 2, 5, 3

- #of page faults = 7 (3 cold-start)
- LRU shown to be useful for workloads with access locality
- Implementation of LRU using the accessed-bit is not easy, approximated using CLOCK